



DEEP LEARNING



Presentation - AE401A

- Aditya Raghuwanshi (170052)

Introduction




- Deep Learning uses a “Neural Networks” architecture to build models.
- NN architecture is designed to work the same way a **human brain** does.
- Deep Learning, unlike other machine learning techniques, performs **feature engineering**.
- Deep Learning suits our ideal machine learning model imagination, “**dump the data into the model and get prediction**”.



Motivation



- DL outperform other techniques if the data size is large compared to other algos.
- Deep Learning is fueled by “**Big Data Era**” and advancements in computational power.
- It is used by world’s largest companies.



facebook

IBM

Qualcomm





amazon



A Single Neuron (Node)





- It is the fundamental unit of neural network
- It performs a simple **linear regression**, more specifically does 2 operations:
 1. Performs weighted sum of input and a bias term
 2. Maps calculated sum to an activation function
- Each neuron gets **activated** (or fire) by some feature(s) in the input data.
- Takes output of previous layers as input, its output is used as input for next layer.
- All weights and bias terms are the **parameters** to be learned by the model.



Activation Functions



- Theoretically, output of a neuron can vary from $-\infty$ to $+\infty$.
- Act. Fun. maps input to a **desired range** that avoids **overflow** and also outputs values which are more useful.
- **Nonlinear function** are, usually, used to learn complex models
 - It can also learn probabilities of a class
 - They allow **back propagation**
- Linear act. fun. are not useful
- Most widely used activation functions are:
 - Sigmoid
 - Tanh
 - ReLu (Rectified Linear Unit)
 - Leaky ReLu



Building A Neural Networks



- Image shows a simple neural network built with:
 - Input Layer, having 3 nodes
 - A Hidden Layer, having 2 nodes
 - Output Layer with single node

- For layer 1, inputs = $[x_1, x_2, x_3]$

$$a_1 = w_{11}x_1 + w_{12}x_2 + w_{13}x_3$$

$$a_2 = w_{21}x_1 + w_{22}x_2 + w_{23}x_3$$

$$h_1 = f(a_1) \text{ and } h_2 = f(a_2)$$

$$\text{Outputs} = [h_1, h_2]$$


For output layer, inputs = $[h_1, h_2]$

$$b = v_1h_1 + v_2h_2$$

$$\text{Output} = g(b)$$



“f” and “g” are activation functions

- This is known as **Forward Propagation** ➡




Building A “Deep” Neural Networks

- Many neurons are stacked to form a layer and many layers are stacked to form NN.
- Number of neurons in each layer and number of layers (depth) can vary and are **hyperparameter**.
- “Deep” in deep learning actually refers to large “depth” of neural architecture.



Training of a Deep NN


- **Stochastic Gradient Descent** is an optimizing algorithm that adjusts the weights to minimize the loss. It is iterative algorithms which performs steps.
 - Make predictions on some data (batch)
 - Compute Loss on the predictions
 - Update parameters to minimise the loss
- **Iterate** on these step until loss is below certain threshold value.
- Predictions are made using **Forward Propagation**, where we sweeps all layer starting from input layer to output layer, calculating outputs of all neurons.



— Batch gradient descent
— Mini-batch gradient Descent
— Stochastic gradient descent

Back Propagation in NN

- Loss are computed using a function called **Loss Function**.
- In Back Prop, model sweeps layers in reverse order and find partial derivative of loss function wrt. parameters and update them to minimise the loss function.
- This is step is **heart of Deep Learning** as “**learning**” happens here.
- Learning Rate** and **Batch Size** are hyperparameters needed to be tuned using validation data.





Overfitting

- Deep Neural Networks performs “too well” on not so complex data which drops performance on test data.
- **Solution: Regularization**
 - Weight Decay Method
 - Dropout Method
 - Early Stopping

- **Weight Decay Method**

Total Loss: $L = (\sum_i L_i^2 + \lambda \cdot \sum_i \|w_i\|_F^2) / m$
It brings activation closer to the origin.

So, as area near origin lies in linear range and model **CANNOT** learn complex nonlinear models.




What actually is **happening** inside a NN?


- During training, each neuron learns a specific feature.
- Neurons in shallow layer learns simple feature (lines or curves) whereas in deeper layer learns complex features (facial features like nose).
- Each neuron in a single layer learns features of same complexity (nearly).
- It is necessary to **randomly initialise** weights to a non-zero value, otherwise all neurons in one single layer will always same value (learn same feature).



Visualizing features learned at each layer




Image



"Audi A7"

Visualizing a Working NN!



In each layer neurons are activated when feature they learned is detected in the image.

Disadvantages of Deep Learning



- Large dataset are required to train a “good” deep neural network.
- Using “Sigmoid” Act. Fun., training slows down at the end.
 - Solution: Use **ReLU** act. fun.
- A good NN model contain thousands or even millions of parameters.
- Have to write large complex code to get a good model.
 - Solution: Use frameworks like **TensorFlow**, **Keras**, etc.
- Training process requires high computational power and is time taking.
 - Solution: **Transfer Learning**
- Require to tune lots of hyperparameters
 - Learning Rate
 - Number of Layers
 - Number of neuron in each layer
 - λ (in regularization)
 - Number of iterations (epoch)
 - Batch Size

References



- Kaggle DL course: <https://www.kaggle.com/learn/intro-to-deep-learning>
- DL Explained (Melanie Swan): <https://www.slideshare.net/lablogga/deep-learning-explained>
- CS771A: Intro to Machine Learning Slides, IIT Kanpur by Prof Purushottam Kar
- Deep Learning Coursera, by Andrew NG:
https://www.coursera.org/programs/indian-institute-of-technology-kanpur-on-coursera-4adct?collectionId=¤tTab=MY_COURSES&productId=W62RsyrdEeeFQQqyuQaohA&productType=s12n&showMiniModal=true
- 3Blue1Brown: <https://www.youtube.com/watch?v=aircAruvnKk&t=108s>
- <https://medium.com/@xzz201920/activation-functions-linear-non-linear-in-deep-learning-relu-sigmoid-softmax-swish-leaky-relu-a6333be712ea>
- <https://medium.com/@jorgesleonel/hyperparameters-in-machine-deep-learning-ca69ad10b981>
- <https://towardsdatascience.com/why-deep-learning-is-needed-over-traditional-machine-learning-1b6a99177063#:~:text=When%20there%20is%20lack%20of,language%20processing%2C%20and%20speech%20recognition.>
- <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>
- <https://www.cleanpng.com/png-rubbish-bins-waste-paper-baskets-hazardous-waste-w-4062584/>
- Google Images



Questions???

Some food for thoughts:

- **Combination of traditional machine learning techniques with DL!**
- **To decrease training time, can we use some pre-computational methods?**
 - **What NN architecture to use for text and image data?**