

# Optimus III

CS771: Introduction to Machine Learning

Purushottam Kar

# Recap of Last Lecture

2

Notion of *subgradient* that elegantly extends the notion of derivative to non-differentiable functions that are nevertheless convex

Subgradient calculus: scaling, sum, chain, max rules

Using the first order optimality condition to solve simple optimization problems

Using (sub)gradient descent to solve general optimization problems

Notions of *initialization*, *step length*, *convergence*



# Behind the scenes in GD

3

$$f(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|_2^2 + C \cdot \sum_{i=1}^n [1 - y^i \cdot \mathbf{w}^\top \mathbf{x}^i]_+ \text{ (ignore bias } b \text{ for now)}$$

$$\nabla f(\mathbf{w}) = \mathbf{w} + C \cdot \sum_{i=1}^n g^i y^i \cdot \mathbf{x}^i, \text{ where } g^i \in \nabla \ell_{\text{hinge}}(y^i \cdot \mathbf{w}^\top \mathbf{x}^i)$$

$$\mathbf{w}^{\text{new}} = \mathbf{w} - \eta \cdot \nabla f(\mathbf{w}) = (1 - \eta) \cdot \mathbf{w} - \eta C \cdot \sum_{i=1}^n g^i y^i \cdot \mathbf{x}^i$$

Assume  $n = 1$  for a moment for sake of understanding

$$\mathbf{w}^{\text{new}} = (1 - \eta) \cdot \mathbf{w} - \eta C \cdot g^1 y^1 \cdot \mathbf{x}^1$$

*Small  $\eta$ :  $(1 - \eta)$  is large  $\Rightarrow$  do not change  $\mathbf{w}$  too much!*

*Large  $\eta$ : Feel free to change  $\mathbf{w}$  as much as the gradient dictates*

*If  $\mathbf{w}$  does well on  $(\mathbf{x}^1, y^1)$ , say  $y^1 \cdot \mathbf{w}^\top \mathbf{x}^1 > 1$ , then  $g^1 = 0$*

*If  $\mathbf{w}$  does badly on  $(\mathbf{x}^1, y^1)$ , say  $y^1 \cdot \mathbf{w}^\top \mathbf{x}^1 < 0$ , then  $g^1 = -1$*


$$\mathbf{w}^{\text{new}} = (1 - \eta) \cdot \mathbf{w} + \eta C \cdot y^1 \cdot \mathbf{x}^1$$

$$y^1 \cdot (\mathbf{w}^{\text{new}})^\top \cdot \mathbf{x}^1 = (1 - \eta) y^1 \cdot \mathbf{w}^\top \mathbf{x}^1 + \eta C \cdot \|\mathbf{x}^1\|_2^2$$



# Behind the scene

So gradient descent, although a mathematical tool from calculus, actually tries very actively to make the model perform better on all data points



$$f(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|_2^2 + C \cdot \sum_{i=1}^n \max(0, 1 - y^i \cdot \mathbf{w}^\top \mathbf{x}^i)$$

$$\nabla f(\mathbf{w}) = \mathbf{w} + C \cdot \sum_{i=1}^n g^i y^i \cdot \mathbf{x}^i, \text{ where } g^i \in \nabla \ell_{\text{hinge}}(y^i \cdot \mathbf{w}^\top \mathbf{x}^i)$$

$$\mathbf{w}^{\text{new}} = \mathbf{w} - \eta \cdot \nabla f(\mathbf{w}) = (1 - \eta) \cdot \mathbf{w} - \eta C \cdot \sum_{i=1}^n g^i y^i \cdot \mathbf{x}^i$$

Assume  $n = 1$  for a moment for sake of understanding

$$\mathbf{w}^{\text{new}} = (1 - \eta) \cdot \mathbf{w} - \eta C \cdot g^1 y^1 \cdot \mathbf{x}^1$$

*Small  $\eta$ :  $(1 - \eta)$  is large  $\Rightarrow$  do not change  $\mathbf{w}$  too much*

*Large  $\eta$ : Feel free to change  $\mathbf{w}$  as much as the gradient dictates*

*If  $\mathbf{w}$  does well on  $(\mathbf{x}^1, y^1)$ , say  $y^1 \cdot \mathbf{w}^\top \mathbf{x}^1 > 1$ , then  $g^1 = 0$*

*If  $\mathbf{w}$  does badly on  $(\mathbf{x}^1, y^1)$ , say  $y^1 \cdot \mathbf{w}^\top \mathbf{x}^1 < 0$ , then  $g^1 = 1$*

$$\mathbf{w}^{\text{new}} = (1 - \eta) \cdot \mathbf{w} + \eta C \cdot y^1 \cdot \mathbf{x}^1$$

$$y^1 \cdot (\mathbf{w}^{\text{new}})^\top \cdot \mathbf{x}^1 = (1 - \eta) y^1 \cdot \mathbf{w}^\top \mathbf{x}^1 + \eta C \cdot \|\mathbf{x}^1\|_2^2$$

No change to  $\mathbf{w}$  due to the data point  $(\mathbf{x}^1, y^1)$

$\mathbf{w}^{\text{new}}$  may get much better margin on  $(\mathbf{x}^1, y^1)$  than  $\mathbf{w}$

# Stochastic Gradient Method

5

$\nabla f(\mathbf{w}) = \mathbf{w} + C \cdot \sum_{i=1}^n g^i y^i \cdot \mathbf{x}^i$ , where  $g^i \in \nabla \ell_{\text{hinge}}(y^i \cdot \mathbf{w}^\top \mathbf{x}^i)$


Calculating each  $g^i$  takes  $\mathcal{O}(d)$  time since  $\mathbf{w}, \mathbf{x}^i \in \mathbb{R}^d$  - total  $\mathcal{O}(nd)$

At each time, choose a random data point  $(\mathbf{x}^{i_t}, y^{i_t})$

$\nabla f(\mathbf{w}) \approx \mathbf{w} + C \cdot g^{i_t} y^{i_t} \cdot \mathbf{x}^{i_t}$  - only  $\mathcal{O}(d)$  time!!

**Warning:** may have to perform several SGD steps than we had to do with GD but each SGD step is much cheaper than a GD step

We take a random data point to avoid being unlucky (also it is cheap)



Do we really need to spend so much time on just one update?

Especially in the beginning, when we are far away from the optimum!

Initially, all we need is a general direction in which to move

No, SGD gives a cheaper way to perform gradient descent



# Mini-batch SGD

6

If data is very diverse, the “stochastic” gradient may vary quite a lot depending on which random data point is chosen

This is called *variance* (more on this later) but this can slow down the SGD process – make it jittery


One solution, choose more than one random point

At each step, choose  $B$  random data points ( $B = \text{mini batch size}$ ) without replacement, say  $(x^{i_1}_t, y^{i_1}_t), \dots, (x^{i_B}_t, y^{i_B}_t)$  and use

$$\nabla f(\mathbf{w}) \approx \mathbf{w} + C \cdot \sum_{b=1}^B g^{i_b}_t y^{i_b}_t \cdot \mathbf{x}^{i_b}_t$$

Takes  $\mathcal{O}(Bd)$  time to execute MBSGD – more expensive than SGD

Notice that if  $B = n$  then MBSGD becomes plain GD



# Constrained Optimization

7

Recall that an optimization problem has an objective and constraints

$\min_x f(x)$

Objective

Constraints

such that  $p(x) < 0$   
and  $q(x) > 0$  etc.

The set of points that satisfy *all* the constraints is called the *feasible set*  $\mathcal{C}$   
 $\mathcal{C} \triangleq \{x : p(x) < 0 \text{ and } q(x) > 0 \text{ and } \dots\}$

Problems with constraints more challenging

## Method 1: Interior Point Method

*Find a way to initialize within  $\mathcal{C}$  and then take steps that never go out*

*A very powerful family of methods – also very involved*

*Not extremely popular in machine learning – can be expensive*

*Beyond the scope of CS771*



# Constrained Optimization

8


## Method 2: Projected Gradient Descent

*Perform (stochastic) gradient descent as usual. However, if this causes us to step outside the feasible set  $\mathcal{C}$ , go back to the feasible set*

*Process of “going back” into  $\mathcal{C}$ : projection step*

$$\hat{\mathbf{x}} = \Pi_{\mathcal{C}}(\mathbf{x}) = \arg \min_{\mathbf{z} \in \mathcal{C}} \|\mathbf{x} - \mathbf{z}\|_2^2$$

*Warning: works only if  $\mathcal{C}$  is such that projection step is easy*



### PROJECTED (SUB)GRADIENT DESCENT

1. Choose  $\mathbf{g}^t \in \partial f(\mathbf{w}^t)$ , step length  $\eta_t$
2. Update  $\mathbf{u}^{t+1} \leftarrow \mathbf{w}^t - \eta_t \cdot \mathbf{g}^t$
3. Project  $\mathbf{w}^{t+1} \leftarrow \Pi_{\mathcal{C}}(\mathbf{u}^{t+1})$
4. Repeat until convergence






# A few useful Projections

$$\hat{\mathbf{x}} = \Pi_{\mathcal{C}}(\mathbf{x})$$


9

$$\mathcal{C} = \{\mathbf{x} : \|\mathbf{x}\|_2 \leq 1\}$$



$$\hat{\mathbf{x}} = \begin{cases} \mathbf{x} & \text{if } \|\mathbf{x}\|_2 \leq 1 \\ \frac{\mathbf{x}}{\|\mathbf{x}\|_2} & \text{if } \|\mathbf{x}\|_2 > 1 \end{cases}$$

$$\mathcal{C} = \{\mathbf{x} : \mathbf{x}_i \geq 0\}$$



$$\hat{\mathbf{x}}_i = \begin{cases} \mathbf{x}_i & \text{if } \mathbf{x}_i \geq 0 \\ 0 & \text{if } \mathbf{x}_i < 0 \end{cases}$$



# Constrained Optimization

10

## Method 3: Creating a Dual Problem

*Suppose we wish to solve*

$$\begin{aligned} \min_{\mathbf{x} \in \mathbb{R}^d} & f(\mathbf{x}) \\ \text{s.t.} & g(\mathbf{x}) \leq 0 \end{aligned}$$

*Trick: sneak this constraint into the objective*


*Construct a barrier function  $r(\mathbf{x})$  so that  $r(\mathbf{x}) = 0$  if  $g(\mathbf{x}) \leq 0$  and  $r(\mathbf{x}) = \infty$  otherwise, and simply solve*

*Easy to see that both problems have the same solution*

*One very elegant way to construct such a barrier is the following*

$$r(\mathbf{x}) = \max_{\alpha \geq 0} \alpha \cdot g(\mathbf{x})$$

Thus, we want to solve  $\min_{\mathbf{x} \in \mathbb{R}^d} \left\{ f(\mathbf{x}) + \max_{\alpha \geq 0} \alpha \cdot g(\mathbf{x}) \right\}$



# Constrained Optimization

## Method 3: Creating a Dual Problem

Suppose we wish to solve

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) + r(\mathbf{x})$$

*Trick: sneak this constraint into the objective*


Construct a barrier function  $r(\mathbf{x})$  so that  $r(\mathbf{x}) = 0$  if  $g(\mathbf{x}) \leq 0$  and  $r(\mathbf{x}) = \infty$  otherwise, and simply solve

Easy to see that both problems have the same solution

One very elegant way to construct such a barrier is the following

$$r(\mathbf{x}) = \max_{\alpha \geq 0} \alpha \cdot g(\mathbf{x})$$

Thus, we want to solve  $\min_{\mathbf{x} \in \mathbb{R}^d} \left\{ f(\mathbf{x}) + \max_{\alpha \geq 0} \alpha \cdot g(\mathbf{x}) \right\}$



# Constrained Optimization

## Method 3: Creating a Dual Problem

Suppose we wish to solve

$$\min_{\mathbf{x} \in \mathbb{R}^d} f(\mathbf{x}) + r(\mathbf{x})$$

*Trick: sneak this constraint into the objective*

Construct a barrier function  $r(\mathbf{x})$  so that  $r(\mathbf{x}) = 0$


Let us see how to handle multiple constraints and equality constraints

Hmm ... we still have a constraint here, but a very simple one i.e.  $\alpha \geq 0$

One very elegant way to construct such a barrier is the following

$$\text{Same as } \min_{\mathbf{x} \in \mathbb{R}^d} \left\{ \max_{\alpha \geq 0} \{f(\mathbf{x}) + \alpha \cdot g(\mathbf{x})\} \right\}$$

Thus, we want to solve



# A few Cleanup Steps

13

**Step 1:** Convert your problem to a minimization problem

$$\max f(\mathbf{x}) \rightarrow \min -f(\mathbf{x})$$

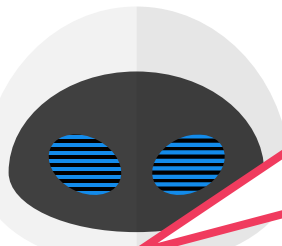
**Step 2:** Convert all inequality constraints to  $\leq$  constraints

$$g(\mathbf{x}) \geq 0 \rightarrow -g(\mathbf{x}) \leq 0$$

**Step 3:** Convert all equality constraints to two inequality constraints


$$s(\mathbf{x}) = 0 \rightarrow s(\mathbf{x}) \leq 0, -s(\mathbf{x}) \leq 0$$

**Step 4:** For each constraint we now have, introduce a new variable  
e.g. if we have  $C$  inequality constraints  $g_1(\mathbf{x}) \leq 0, \dots, g_C(\mathbf{x}) \leq 0$ ,  
introduce  $C$  new variables  $\alpha_1, \dots, \alpha_C$



The variables of the original optimization problem, e.g.  $\mathbf{x}$  in this case, are called the *primal variables* by comparison

These new variables are called *dual variables* or sometimes even called *Lagrange multipliers*



# The Lagrangian

14

$$\begin{array}{ll}\min_{\mathbf{x}} & f(\mathbf{x}) \\ \text{s.t.} & g_1(\mathbf{x}) \leq 0 \\ & g_2(\mathbf{x}) \leq 0 \\ & \vdots \\ & g_C(\mathbf{x}) \leq 0\end{array}$$

$\mathcal{L}(\mathbf{x}, \boldsymbol{\alpha}) = f(\mathbf{x}) + \sum_{c=1}^C \boldsymbol{\alpha}_c \cdot g_c(\mathbf{x})$  called the *Lagrangian* of the problem

If  $\mathbf{x}$  violates even one constraint, we have

$$\max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \boldsymbol{\alpha}_c \geq 0}} \{\mathcal{L}(\mathbf{x}, \boldsymbol{\alpha})\} = \infty$$

If  $\mathbf{x}$  satisfies every single constraint, we have

$$\max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \boldsymbol{\alpha}_c \geq 0}} \{\mathcal{L}(\mathbf{x}, \boldsymbol{\alpha})\} = f(\mathbf{x})$$

$$\min_{\mathbf{x} \in \mathbb{R}^d} \left\{ \max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \boldsymbol{\alpha}_c \geq 0}} \left\{ f(\mathbf{x}) + \sum_{c=1}^C \boldsymbol{\alpha}_c \cdot g_c(\mathbf{x}) \right\} \right\}$$



# The Lagrangian

15

$$\begin{array}{ll}\min_{\mathbf{x}} & f(\mathbf{x}) \\ \text{s.t.} & g_1(\mathbf{x}) \leq 0 \\ & g_2(\mathbf{x}) \leq 0 \\ & \vdots \\ & g_C(\mathbf{x}) \leq 0\end{array}$$

$\mathcal{L}(\mathbf{x}, \boldsymbol{\alpha}) = f(\mathbf{x}) + \sum_{c=1}^C \boldsymbol{\alpha}_c \cdot g_c(\mathbf{x})$  called the *Lagrangian* of the problem


If  $\mathbf{x}$  violates even one constraint, we have

$$\max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \boldsymbol{\alpha}_c \geq 0}} \{\mathcal{L}(\mathbf{x}, \boldsymbol{\alpha})\} = \infty$$

If  $\mathbf{x}$  satisfies every single constraint, we have

This is just a nice way of rewriting the above problem

$$\min_{\mathbf{x} \in \mathbb{R}^d} \left\{ \max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \boldsymbol{\alpha}_c \geq 0}} \left\{ f(\mathbf{x}) + \sum_{c=1}^C \boldsymbol{\alpha}_c \cdot g_c(\mathbf{x}) \right\} \right\}$$



# The Dual Problem

16

The original optimization problem is also called the *primal problem*

Recall: variables of the original problem e.g.  $\mathbf{x}$  called *primal variables*

Using the Lagrangian, we rewrote the primal problem as

$$\min_{\mathbf{x} \in \mathbb{R}^d} \left\{ \max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \alpha_c \geq 0}} \left\{ f(\mathbf{x}) + \sum_{c=1}^C \alpha_c \cdot g_c(\mathbf{x}) \right\} \right\}$$

The dual problem is obtained by simply switching order of min/max

$$\max_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \alpha_c \geq 0}} \left\{ \min_{\mathbf{x} \in \mathbb{R}^d} \left\{ f(\mathbf{x}) + \sum_{c=1}^C \alpha_c \cdot g_c(\mathbf{x}) \right\} \right\}$$

In some cases, the dual problem is easier to solve than the primal





Let  $\hat{\mathbf{x}}^P, \hat{\boldsymbol{\alpha}}^P$  be the solutions to the primal problem i.e.

$$(\hat{\mathbf{x}}^P, \hat{\boldsymbol{\alpha}}^P) = \operatorname{argmin}_{\mathbf{x} \in \mathbb{R}^d} \left\{ \operatorname{argmax}_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \alpha_c \geq 0}} \left\{ f(\mathbf{x}) + \sum_{c=1}^C \alpha_c \cdot g_c(\mathbf{x}) \right\} \right\}$$

Let  $\hat{\mathbf{x}}^D, \hat{\boldsymbol{\alpha}}^D$  be the solutions to the dual problem i.e.

$$(\hat{\mathbf{x}}^D, \hat{\boldsymbol{\alpha}}^D) = \operatorname{argmax}_{\substack{\boldsymbol{\alpha} \in \mathbb{R}^C \\ \alpha_c \geq 0}} \left\{ \operatorname{argmin}_{\mathbf{x} \in \mathbb{R}^d} \left\{ f(\mathbf{x}) + \sum_{c=1}^C \alpha_c \cdot g_c(\mathbf{x}) \right\} \right\}$$

**Strong Duality:**  $\hat{\mathbf{x}}^P = \hat{\mathbf{x}}^D$  if the original problem is convex and “nice”

**Complementary Slackness:**  $\hat{\alpha}_c^D \cdot g_c(\hat{\mathbf{x}}^D) = 0$  for all constraints  $c$

**Note:** not complimentary but complementary 😊



# Hard SVM without a bias

18

$\min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{2} \|\mathbf{w}\|_2^2$  such that  $1 - y^i \cdot \mathbf{w}^\top \mathbf{x}^i \leq 0$  for all  $i \in [n]$

$n$  constraints so we need  $n$  dual variables i.e.  $\boldsymbol{\alpha} \in \mathbb{R}^n$

Lagrangian:  $\mathcal{L}(\mathbf{w}, \boldsymbol{\alpha}) = \frac{1}{2} \|\mathbf{w}\|_2^2 + \sum_{i=1}^n \alpha_i (1 - y^i \cdot \mathbf{w}^\top \mathbf{x}^i)$

Primal problem:  $\operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \left\{ \operatorname{argmax}_{\boldsymbol{\alpha} \geq 0} \left\{ \frac{1}{2} \|\mathbf{w}\|_2^2 + \sum_{i=1}^n \alpha_i (1 - y^i \mathbf{w}^\top \mathbf{x}^i) \right\} \right\}$

Dual problem:  $\operatorname{argmax}_{\boldsymbol{\alpha} \geq 0} \left\{ \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \left\{ \frac{1}{2} \|\mathbf{w}\|_2^2 + \sum_{i=1}^n \alpha_i (1 - y^i \mathbf{w}^\top \mathbf{x}^i) \right\} \right\}$

The dual problem can be greatly simplified!



# Simplifying the Dual Problem

19

Note that the inner problem in the dual problem is

$$\operatorname{argmax}_{\alpha \geq 0} \left\{ \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \left\{ \frac{1}{2} \|\mathbf{w}\|_2^2 + \sum_{i=1}^n \alpha_i (1 - y^i \mathbf{w}^\top \mathbf{x}^i) \right\} \right\}$$

Since this is an unconstrained problem with a convex and differentiable objective, we can apply first order optimality to solve it completely ☺

If we set the gradient to zero, we will get  $\mathbf{w} = \sum_{i=1}^n \alpha_i y^i \cdot \mathbf{x}^i$

Substituting this back in the dual problem we get


$$\operatorname{argmax}_{\alpha \geq 0} \left\{ \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i, \mathbf{x}^j \rangle \right\}$$

This is actually the problem several solvers (e.g. libsvm, sklearn) solve



# Simplifying the Dual Problem

Once you get optimal values of  $\alpha$ , use  $\mathbf{w} = \sum_{i=1}^n \alpha_i y^i \cdot \mathbf{x}^i$  to get optimal value of  $\mathbf{w}$



$$\operatorname{argmax}_{\alpha \geq 0} \left\{ \operatorname{argmin}_{\mathbf{w} \in \mathbb{R}^d} \left\{ \frac{1}{2} \|\mathbf{w}\|_2^2 + \sum_{i=1}^n \alpha_i (1 - y^i \mathbf{w}^\top \mathbf{x}^i) \right\} \right\}$$

Since this is an unconstrained problem with a convex and differentiable objective, we can apply first order optimality to solve it completely ☺

If we set the gradient to zero, we will get  $\mathbf{w} = \sum_{i=1}^n \alpha_i y^i \cdot \mathbf{x}^i$

Substituting this back in the dual problem we get

$$\operatorname{argmax}_{\alpha \geq 0} \left\{ \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y^i y^j \langle \mathbf{x}^i, \mathbf{x}^j \rangle \right\}$$

This is actually the problem several solvers (e.g. libsvm, sklearn) solve



# Support Vectors


Recall: we have  $\alpha_i$  for every data point

After solving the dual problem, the data points for which  $\alpha_i \neq 0$ : **Support Vectors**

Usually we have  $\ll n$  support vectors

Recall: complementary slackness tells us that  $\alpha_i(1 - y^i \mathbf{w}^T \mathbf{x}^i) = 0$  i.e. only those data points can become SVs for which  $y^i \mathbf{w}^T \mathbf{x}^i = 1$  i.e. at margin

The reason these are called *support* vectors has to do with a mechanical interpretation of these objects – need to look at CSVM to understand that



# Support Vectors

Recall: we have  $\alpha_i$  for every data point


After solving the dual problem, the data points for which  $\alpha_i \neq 0$ : **Support Vectors**

Usually we have  $\ll n$  support vectors

Recall: complementary slackness tells us that  $\alpha_i(1 - y^i \mathbf{w}^T \mathbf{x}^i) = 0$  i.e. only those data points can become SVs for which  $y^i \mathbf{w}^T \mathbf{x}^i = 1$  i.e. at margin

The reason these are called *support* vectors has to do with a mechanical interpretation of these objects – need to look at CSVM to understand that

Support Vectors!!



Similar calculations show that if we have a bias term  $b$  as well as slack variables, then the dual looks like

$$\begin{aligned} & \underset{\boldsymbol{\alpha} \in \mathbb{R}^n}{\operatorname{argmax}} \left\{ \sum_{i=1}^n \boldsymbol{\alpha}_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \boldsymbol{\alpha}_i \boldsymbol{\alpha}_j y^i y^j \langle \mathbf{x}^i, \mathbf{x}^j \rangle \right\} \\ & \text{s.t. } \boldsymbol{\alpha}_i \in [0, C], \text{ and } \sum_{i=1}^n \boldsymbol{\alpha}_i y^i = 0 \end{aligned}$$

If we think of  $\boldsymbol{\alpha}^i$  as the “force” that data point  $i$  applies on the hyperplane  $\mathbf{w}$  and  $y^i$  as the direction in which that force is applied, then we have total force zero since  $\sum_{i=1}^n \boldsymbol{\alpha}_i y^i = 0$ . Also since we also have  $\mathbf{w} = \sum_{i=1}^n \boldsymbol{\alpha}_i y^i \cdot \mathbf{x}^i$ , we can think of the torque on the hyperplane as being zero as well. Thus, the support vectors *mechanically support* the hyperplane, hence their name 😊

**Next class:** how to solve the dual problem?

