

Principal Component Analysis III

CS771: Introduction to Machine Learning

Purushottam Kar

Announcements

2

No classes next week on account of mid-semester recess

Classes resume on Monday, October 14

Quiz 3 to be held on October 16



Recap of Last Lecture

3

Applications of PCA

Space savings: $\mathcal{O}((n + d) \cdot k)$ vs $\mathcal{O}(nd)$

Time savings (w.r.t linear models): $\mathcal{O}((n + d) \cdot k)$ vs $\mathcal{O}(nd)$

Noise removal: dims with small(er) singular values may be noise

Left singular vectors give us (new) low-dimensional representation of data

Example application: foreground-background separation in videos

Learning prototypes: right singular vecs can be treated as prototypes

Data can be linearly approx. in terms of these prototypes (not so in GMM etc)

Example applications: eigenfaces, LSA, collaborative filtering

Many (equivalent) interpretations of PCA

Gives smallest reconstruction error, largest variance preservation



Probabilistic PCA

4

The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



Probabilistic PCA

4

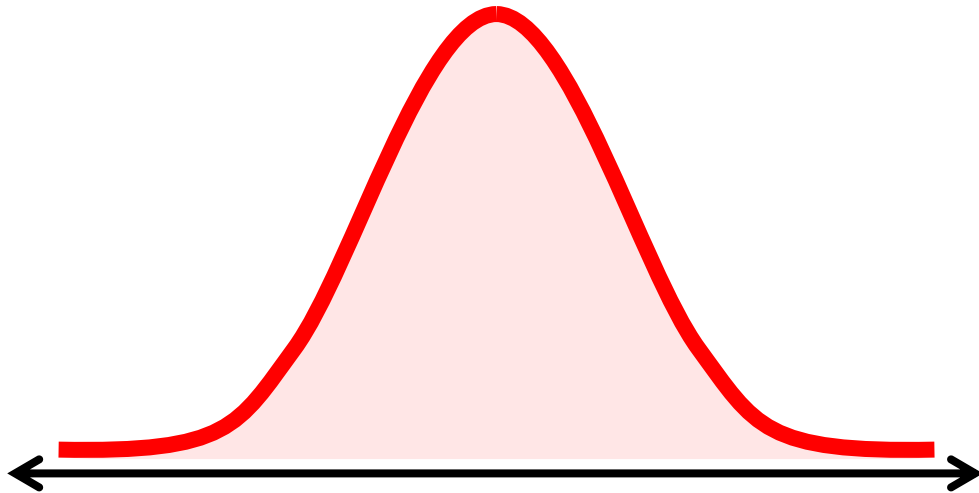
The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



Probabilistic PCA

4

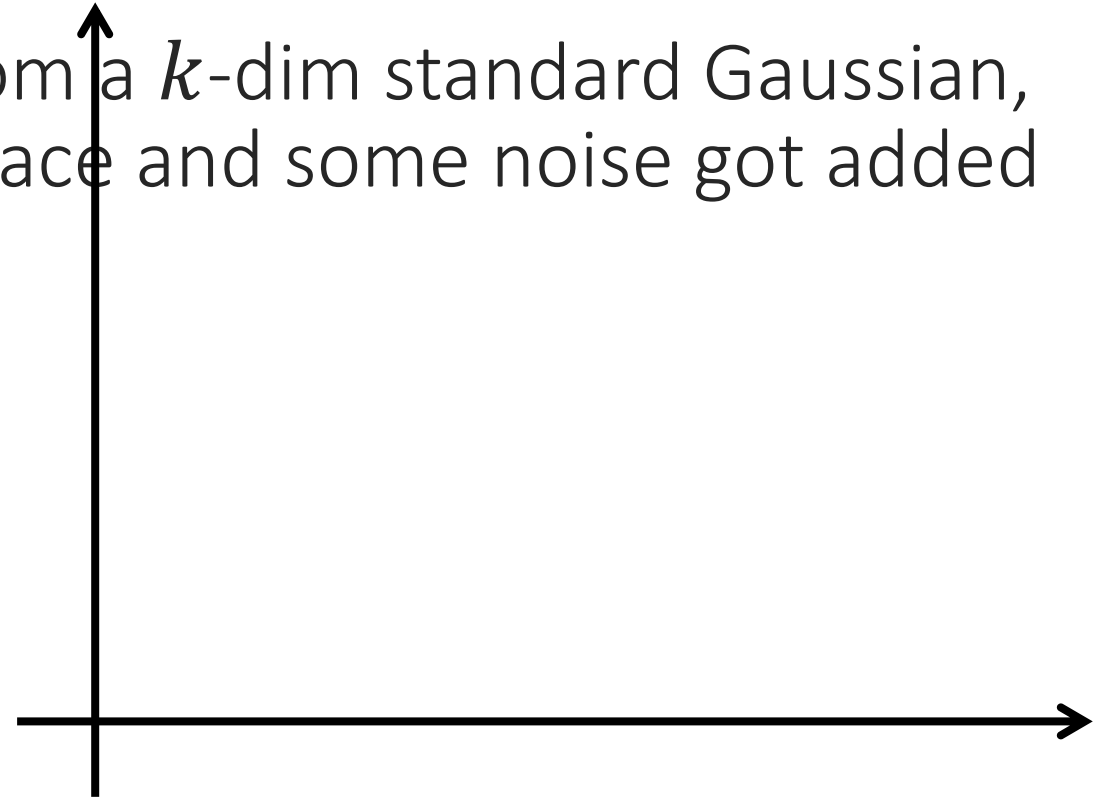
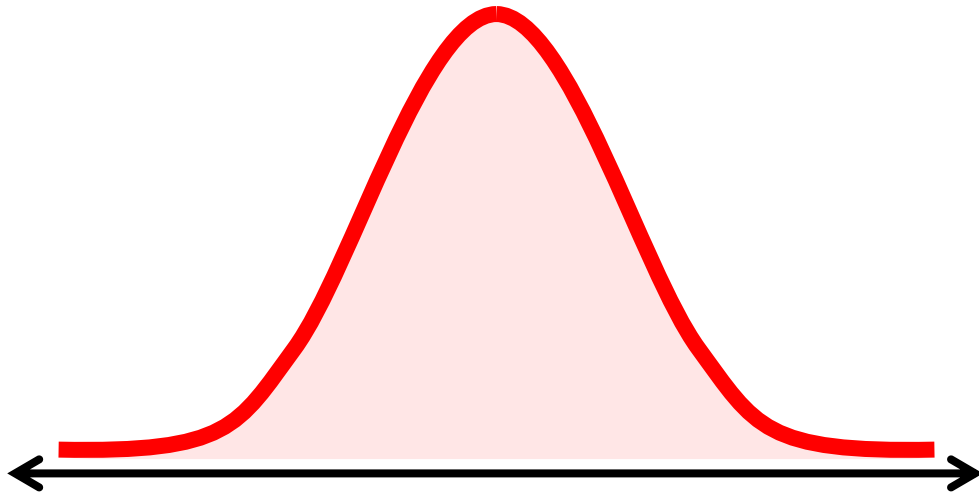
The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



Probabilistic PCA

4

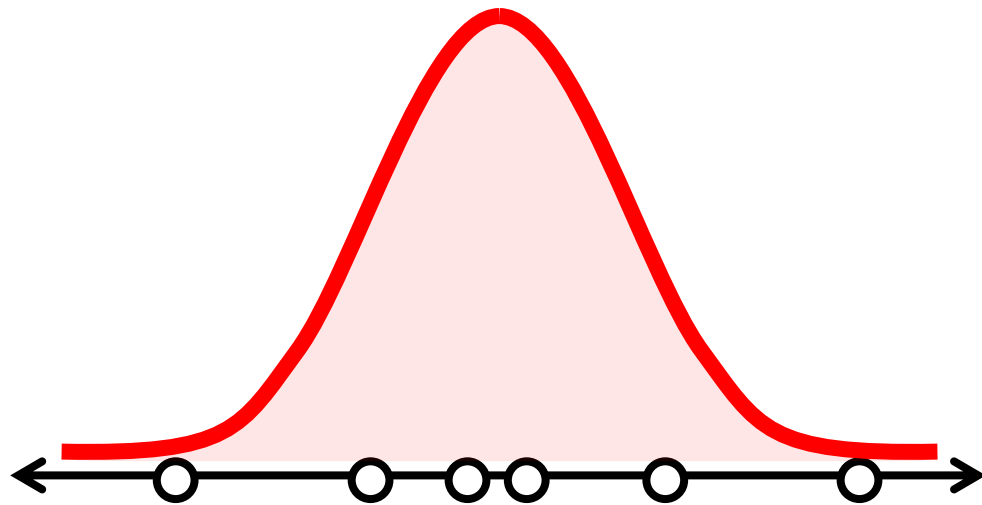
The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



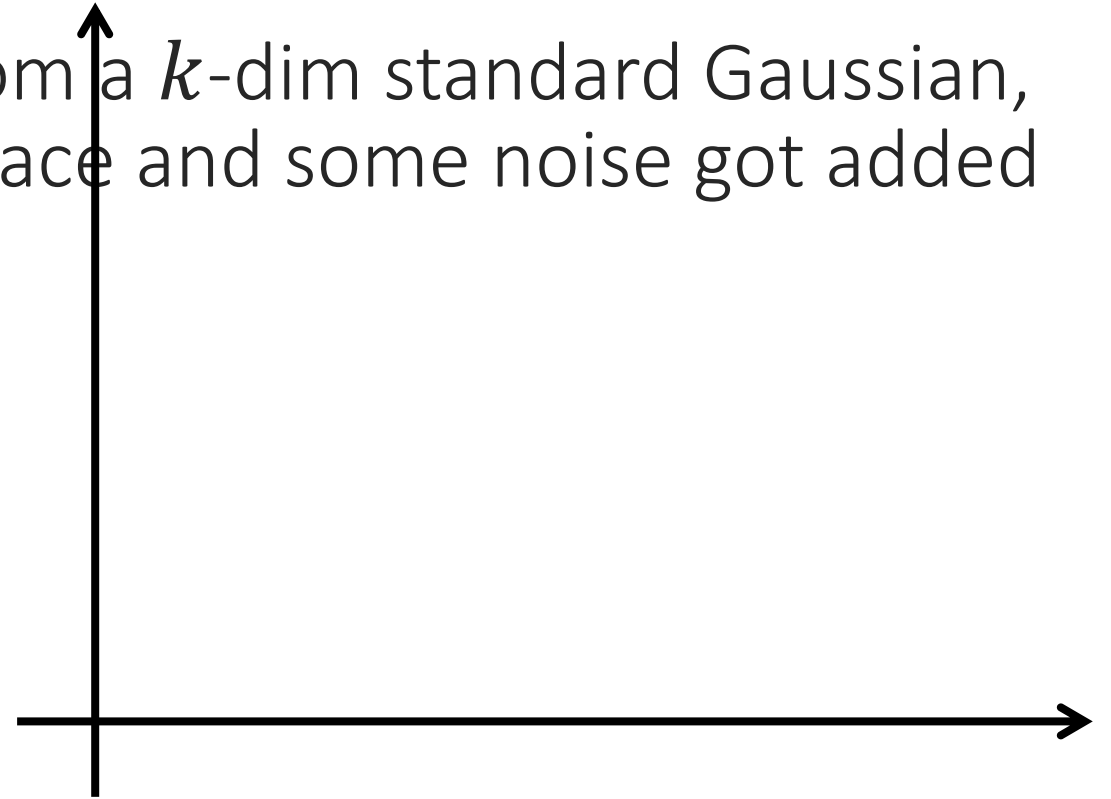
Probabilistic PCA

4

The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added




$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$



Probabilistic PCA

4


The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W \mathbf{z}^i$$

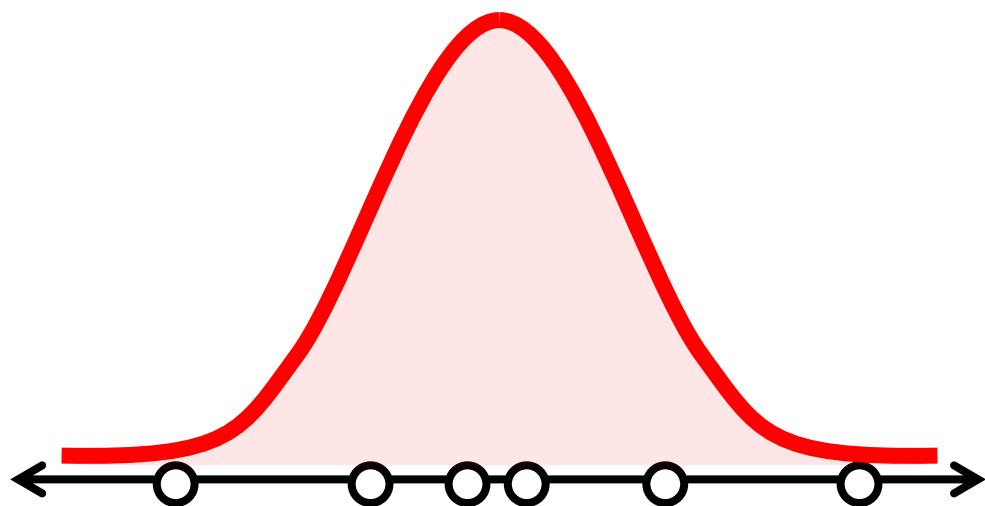
$$W \in \mathbb{R}^{d \times k}$$



Probabilistic PCA

4

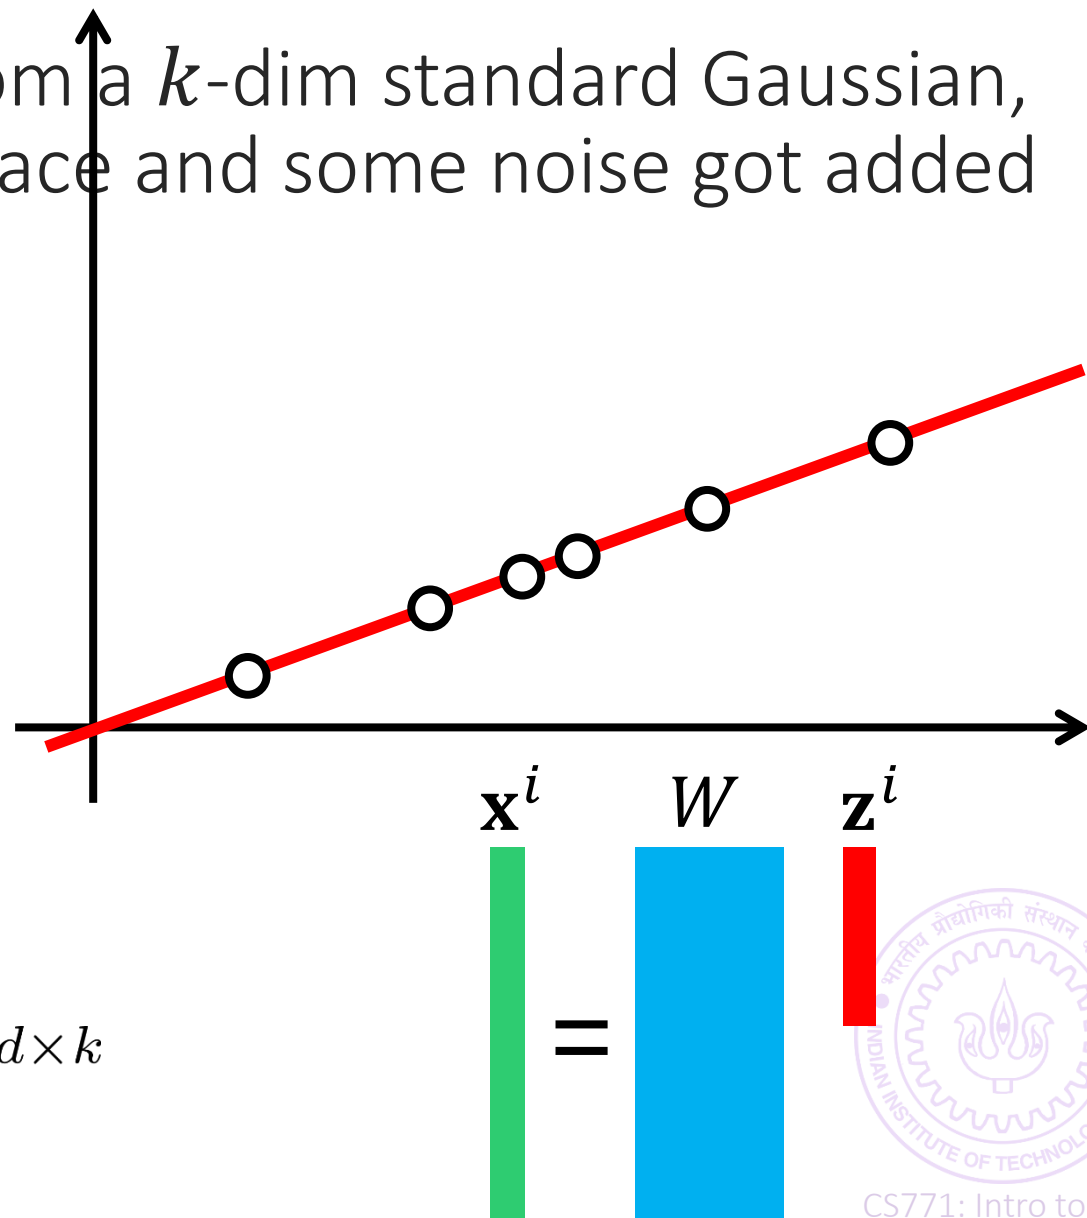
The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W \mathbf{z}^i$$


$$W \in \mathbb{R}^{d \times k}$$



Probabilistic PCA

4


The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W \mathbf{z}^i$$



$$W \in \mathbb{R}^{d \times k}$$



Probabilistic PCA

4

The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added




$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W\mathbf{z}^i + \boldsymbol{\epsilon}^i, \quad \boldsymbol{\epsilon}^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

$$W \in \mathbb{R}^{d \times k}$$

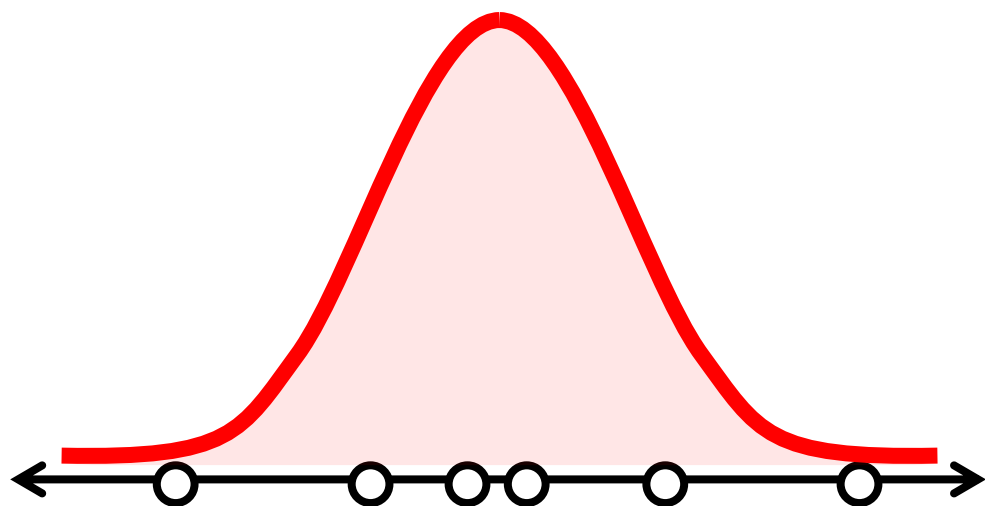
A diagram illustrating the matrix equation $\mathbf{x}^i = W\mathbf{z}^i$. On the left is a tall, thin green vertical bar labeled \mathbf{x}^i . In the middle is a blue horizontal bar labeled W . On the right is a tall, thin red vertical bar labeled \mathbf{z}^i . An equals sign is placed between the green bar and the blue bar, and between the blue bar and the red bar.



Probabilistic PCA

4


The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W\mathbf{z}^i + \boldsymbol{\epsilon}^i, \quad \boldsymbol{\epsilon}^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

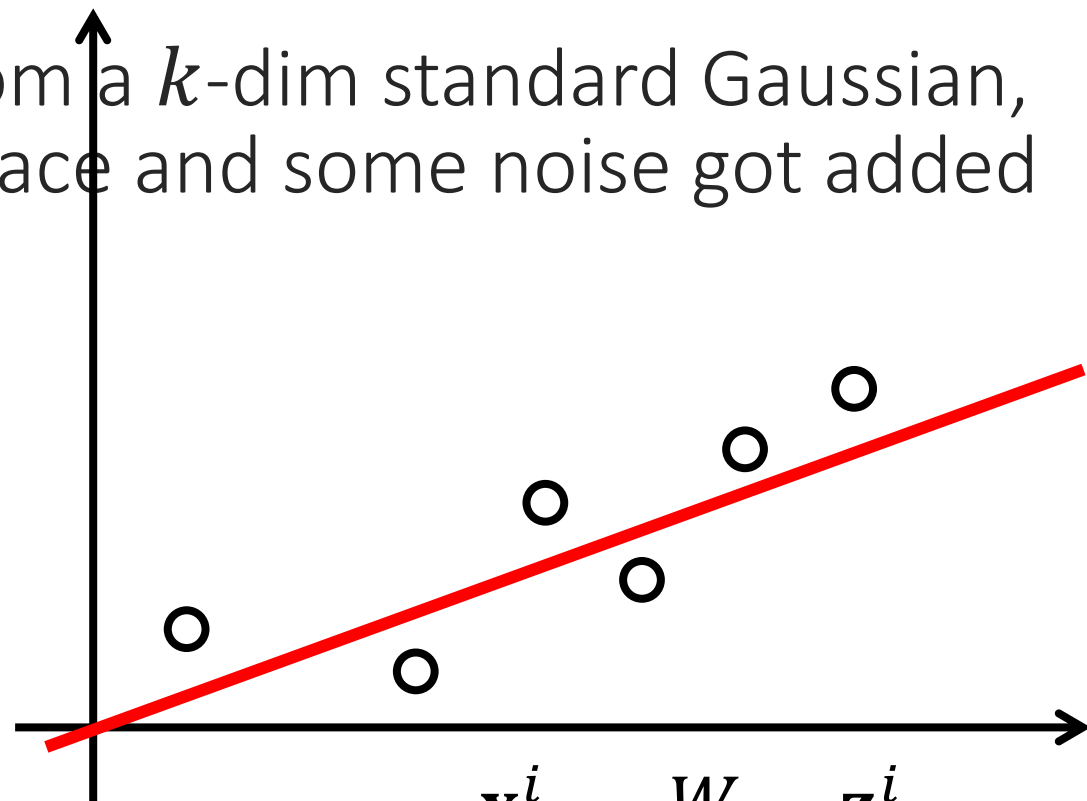

$$W \in \mathbb{R}^{d \times k}$$



Probabilistic PCA

4

The real data was actually sampled from a k -dim standard Gaussian, but got linearly mapped to a d -dim space and some noise got added



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k) \quad \mathbf{x}^i = W \mathbf{z}^i + \epsilon^i, \quad \epsilon^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

Dictionary/Factor
Loading matrix

$$W \in \mathbb{R}^{d \times k}$$

$$\mathbf{x}^i = W \mathbf{z}^i$$



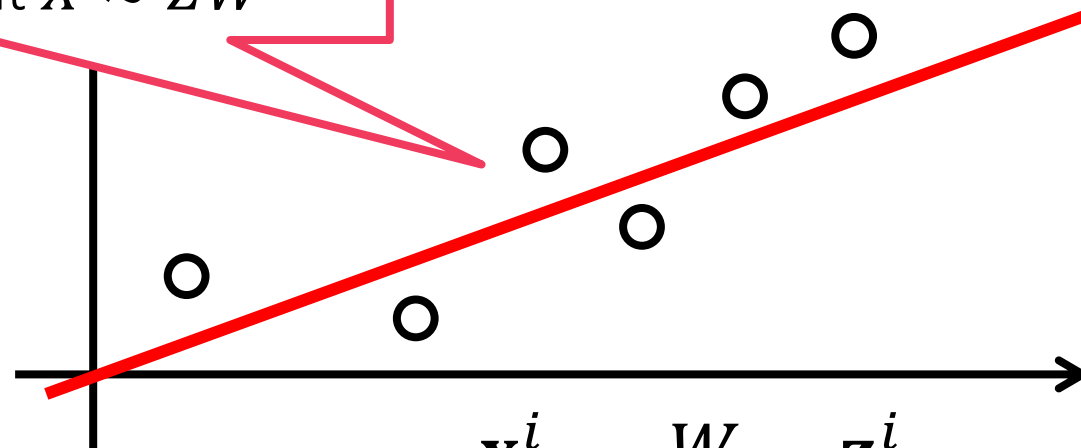

Probabilistic PCA

4

The real data was actually generated from a standard Gaussian, but got linearly mapped

Given $\mathbf{x}^1, \dots, \mathbf{x}^n$, can we recover $W, \mathbf{z}^1, \dots, \mathbf{z}^n$? In other words, given $X \in \mathbb{R}^{n \times d}$, recover $W \in \mathbb{R}^{d \times k}$ and $Z \in \mathbb{R}^{n \times k}$ such that $X \approx ZW^\top$

standard Gaussian, but noise got added



$$\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$$

$$\mathbf{x}^i = W\mathbf{z}^i + \epsilon^i, \quad \epsilon^i \sim \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I_d)$$

Dictionary/Factor
Loading matrix

$$W \in \mathbb{R}^{d \times k}$$

$$\mathbf{x}^i = W \mathbf{z}^i$$



Probabilistic PCA [Tipping and Bishop, 1999]

16

Given samples $X = [\mathbf{x}^1, \dots, \mathbf{x}^n]^\top \in \mathbb{R}^{n \times d}$, we wish to recover W, σ, \mathbf{z}^i

Note: this is a generative problem, i.e. deals with generation of feature vectors

As discussed before, the original data $\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$ are “latent” – not seen

Also clear from the noise model that $\mathbb{P}[\mathbf{x}^i \mid \mathbf{z}^i, \sigma, W] = \mathcal{N}(\mathbf{x}^i; W\mathbf{z}^i, \sigma^2 \cdot I_d)$

More flexible models possible e.g. Factor Analysis – will see later

Will first see how to recover W, σ and then head into recovering \mathbf{z}^i

Some mildly painful integrals later we can get

$$\mathbb{P}[\mathbf{x}^i \mid \sigma, W] = \int_{\mathbf{z} \in \mathbb{R}^k} \mathbb{P}[\mathbf{x}^i \mid \mathbf{z}^i, \sigma, W] \cdot \mathbb{P}[\mathbf{z}^i] d\mathbf{z} = \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I + WW^\top)$$

The above also assumes $\mathbb{P}[\mathbf{z}^i \mid \sigma, W] = \mathbb{P}[\mathbf{z}^i] = \mathcal{N}(\mathbf{z}^i; \mathbf{0}, I_k)$


Thus, to simplify life, we can pretend for a moment that our samples X were really generated from $\mathcal{N}(\mathbf{0}, \sigma^2 \cdot I + WW^\top)$ and there are no \mathbf{z}^i in the picture.

Can we now do MLE to find σ, W ?



Probabilist

It is very unusual for latent variables to simply integrate out like this leaving behind a nice Gaussian density. We got very lucky here 😊. Usually latent variables mean AltOpt/EM



Given samples X

***Note:** this is a generative problem, i.e. deals with generation of feature vectors*
As discussed before, the original data $\mathbf{z}^i \sim \mathcal{N}(\mathbf{0}, I_k)$ are “latent” – not seen
Also clear from the noise model that $\mathbb{P}[\mathbf{x}^i \mid \mathbf{z}^i, \sigma, W] = \mathcal{N}(\mathbf{x}^i; W\mathbf{z}^i, \sigma^2 \cdot I_d)$
More flexible models possible e.g. Factor Analysis – will see later

Will first see how to recover W, σ and then head into recovering \mathbf{z}^i

Some mildly painful integrals later we can get

$$\mathbb{P}[\mathbf{x}^i \mid \sigma, W] = \int_{\mathbf{z} \in \mathbb{R}^k} \mathbb{P}[\mathbf{x}^i \mid \mathbf{z}^i, \sigma, W] \cdot \mathbb{P}[\mathbf{z}^i] d\mathbf{z} = \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I + WW^\top)$$

The above also assumes $\mathbb{P}[\mathbf{z}^i \mid \sigma, W] = \mathbb{P}[\mathbf{z}^i] = \mathcal{N}(\mathbf{z}^i; \mathbf{0}, I_k)$

Thus, to simplify life, we can pretend for a moment that our samples X were really generated from $\mathcal{N}(\mathbf{0}, \sigma^2 \cdot I + WW^\top)$ and there are no \mathbf{z}^i in the picture.

Can we now do MLE to find σ, W ?



An MLE For W, σ in PPCA

18

We get n samples X from $\mathcal{N}(0, \sigma^2 \cdot I + WW^\top)$. The log-likelihood is

$$\log \mathbb{P}[X \mid \sigma, W] = \frac{n}{2} (d \log 2\pi + \log |C| + \text{tr}(C^{-1}S))$$

where $C = WW^\top + \sigma^2 \cdot I_d$ and $S = \frac{1}{n} \cdot X^\top X$

Note: C is always invertible because of $\sigma^2 \cdot I_d$ (imp. since $|WW^\top| = 0$ ☹)

Can apply first order optimality to get MLE (painful derivatives though)

Thankfully, end result is very familiar. Let $S = Q\Lambda Q^\top$ be eigendecomposition of S where $Q = [\mathbf{q}^1, \dots, \mathbf{q}^d]$ and $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$ with $\lambda_1 \geq \lambda_2 \geq \dots$

An ED always exists for S since it is square symmetric

$$\hat{W}_{\text{MLE}} = \hat{Q} \sqrt{\hat{\Lambda} - \hat{\sigma}_{\text{MLE}}^2 \cdot I_k} \text{ where } \hat{Q} = [\mathbf{q}^1, \dots, \mathbf{q}^k], \hat{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_k)$$

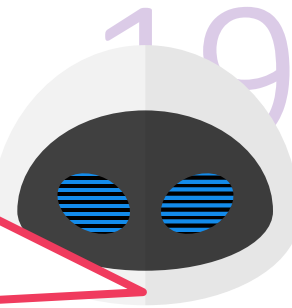
$$\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \lambda_j$$



A

W

If we decide to set $\sigma = 0$ (and not estimate it either) then we get $\hat{W}_{\text{MLE}} = \hat{Q}\sqrt{\hat{\Lambda}}$ which, apart from the scaling with the eigenvalues, is just \hat{V} in PCA! Thus, PCA \equiv PPCA (apart from a scaling factor) under the noiseless assumption 😊



$$\log \mathbb{P}[X \mid \sigma, W] = \frac{n}{2} (d \log 2\pi + \log |C| + \text{tr}(C^{-1}S))$$

$$\text{where } C = WW^T + \sigma^2 \cdot I_d \text{ and } S = \frac{1}{n} \cdot X^T X$$

Note: C is always invertible because of $\sigma^2 \cdot I_d$ (imp. since $|WW^T| = 0$ 😞)

Can apply first order optimality to get MLE (painful derivatives though)

Thankfully, end result is very familiar. Let $S = Q\Lambda Q^T$ be eigendecomposition of S where $Q = [\mathbf{q}^1, \dots, \mathbf{q}^d]$ and $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_d)$ with $\lambda_1 \geq \lambda_2 \geq \dots$

An ED always exists for S since it is square symmetric

$$\hat{W}_{\text{MLE}} = \hat{Q} \sqrt{\hat{\Lambda} - \hat{\sigma}_{\text{MLE}}^2 \cdot I_k} \text{ where } \hat{Q} = [\mathbf{q}^1, \dots, \mathbf{q}^k], \hat{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_k)$$

$$\hat{\sigma}_{\text{MLE}} = \frac{1}{d-k} \sum_{j=k+1}^d \lambda_j$$



We can tweak several moving parts in the PPCA generative process

Can instead assume that $\mathbf{z}^i \sim \mathcal{N}(\boldsymbol{\mu}_z, \Sigma_z)$ and estimate $\boldsymbol{\mu}_z \in \mathbb{R}^k, \Sigma_z \in \mathbb{R}^{k \times k}$

Can assume non-spherical noise $\boldsymbol{\epsilon}^i \sim \mathcal{N}(\mathbf{0}, \Sigma_\epsilon)$ and estimate $\Sigma_\epsilon \in \mathbb{R}^{d \times d}$

A technique called Factor Analysis actually uses non-spherical noise model

Since PPCA is a generative model, it can model missing data too

Suppose we have already found out $\hat{\sigma}_{\text{MLE}}, \hat{W}_{\text{MLE}}$ using clean training data

If test data has missing features $\mathbf{x}^t = [\mathbf{x}_o^t, \mathbf{x}_m^t] \in \mathbb{R}^d$, use fact that marginals of Gaussians are Gaussian. Since we know $\mathbf{x}^t \sim \mathcal{N}(\mathbf{0}, \hat{\sigma}^2 \cdot I + \hat{W}\hat{W}^\top)$, we can see $\mathbb{P}[\mathbf{x}_o^t] = \mathcal{N}(\mathbf{0}, \hat{\sigma}^2 \cdot I_o + \hat{W}_o\hat{W}_o^\top)$ ($\hat{W}_o \in \mathbb{R}^{|\mathcal{O}| \times k}$ has only observed rows)

Missing test data is easier to handle, missing training data more challenging

Need to apply the above trick when defining the likelihood of training data points

Each training data point may have different coordinates missing. If we do not even know when a coordinate has gone missing then this becomes more challenging ☹️

Dimensionality Reduction using PPCA

21

With PCA $X \approx \hat{X} = \hat{U}\hat{\Sigma}\hat{V}^\top$ we got low-dim feat. easily $\tilde{X} = \hat{U}\hat{\Sigma} \in \mathbb{R}^{n \times k}$

These made sense because these are k -dim features which are just a rotation away (using V) from features \hat{X} which we know approximate X very well

With PPCA too we can recover the original low-dim features $\mathbf{z}^i \in \mathbb{R}^k$ by treating them as latent variables and applying AltOpt or EM

Need to be careful since these latent variables are (continuous) vectors now!

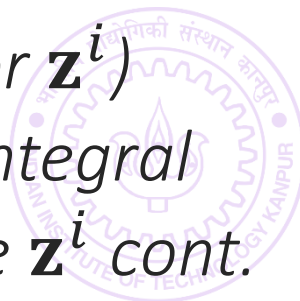
Earlier, we used a shortcut $\mathbb{P}[\mathbf{x}^i \mid \sigma, W] = \mathcal{N}(\mathbf{0}, \sigma^2 \cdot I + WW^\top)$ to get MLE for W, σ . To get hold of \mathbf{z}^i we need proper AltOpt/EM

$$\mathbb{P}[\mathbf{x}^i \mid \sigma, W] = \int_{\mathbb{R}^k} \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] d\mathbf{z}^i$$

AltOpt will approximate integral using a single term (a single value for \mathbf{z}^i)

EM will lower bound the integral using another (easier to compute) integral

Need to replace “sum” over possible values of \mathbf{z}^i with “integral” since \mathbf{z}^i cont.



PPCA – Alternating Optimization

22

We wish to get $\arg \max_{W, \sigma} \sum_{i=1}^n \ln \left(\int_{\mathbb{R}^k} \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] d\mathbf{z}^i \right)$

Same old pesky sum-log-sum (actually sum-log-integral) form – difficult ☹

Approximate integral by its most dominant term $\int_{\mathbb{R}^k} \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] d\mathbf{z}^i \approx \mathbb{P}[\mathbf{x}^i, \tilde{\mathbf{z}}^i \mid \sigma, W]$ with $\tilde{\mathbf{z}}^i = \arg \max_{\mathbf{z}} \mathbb{P}[\mathbf{x}^i, \mathbf{z} \mid \sigma, W] = \arg \max_{\mathbf{z}} \mathbb{P}[\mathbf{z} \mid \mathbf{x}^i, \sigma, W]$

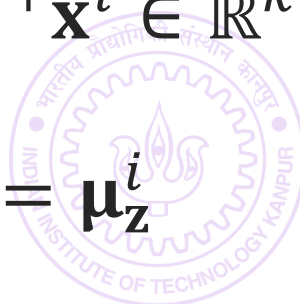
Approx. integral by a single term may not be bad if the dist. $\mathbb{P}[\mathbf{z} \mid \mathbf{x}^i, \sigma, W]$ has small variance – advantage of being cheaper than EM in computation time

Thus, we wish to solve $\arg \max_{W, \sigma, \mathbf{z}^i} \sum_{i=1}^n \ln \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W]$

$\mathbb{P}[\mathbf{z}^i \mid \mathbf{x}^i, \sigma, W] = \mathcal{N}(\mathbf{z}^i; \boldsymbol{\mu}_{\mathbf{z}}^i, \Sigma_{\mathbf{z}})$ where $\boldsymbol{\mu}_{\mathbf{z}}^i = (W^\top W + \sigma^2 \cdot I_k)^{-1} W^\top \mathbf{x}^i \in \mathbb{R}^k$
and $\Sigma_{\mathbf{z}} = \sigma^2 \cdot (W^\top W + \sigma^2 \cdot I_k)^{-1} \in \mathbb{R}^{k \times k}$

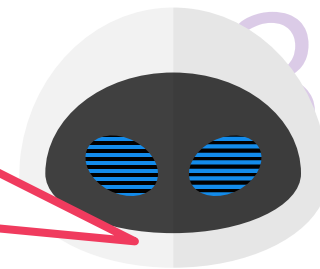
Since for Gaussians, mode is mean, we have $\arg \max_{\mathbf{z}} \mathbb{P}[\mathbf{z} \mid \mathbf{x}^i, \sigma, W] = \boldsymbol{\mu}_{\mathbf{z}}^i$

This gives us one of the alternating updates, let's derive the other



PPCA – A

These derivations are routine but tedious. You can check [BIS] Chapter 12 for details. **Warning:** equation 12.42 in that book has an error. The correct expression is given in this slide



We wish to get $\arg \max_{\mathbf{z}^i} \mathbb{P}(\mathbf{z}^i | \mathbf{x}^i, \sigma, W)$



Because this the mean of \mathbf{z}^i conditioned on \mathbf{x}^i . It is the marginal (unconditional) mean of \mathbf{z}^i that is $\mathbf{0}$. Suppose we know that $\mathbf{x}^i = W\mathbf{z}^i + \epsilon^i$ is a vector far off from the origin. Then it is likely that $W\mathbf{z}^i \neq \mathbf{0}$ as well which immediately tells that $\mathbf{z}^i \neq \mathbf{0}$

We know that $\mathbf{z}^i \sim \mathcal{N}(\mathbf{z}^i; \mathbf{0}, I_k)$. Then how come $\mu_{\mathbf{z}}^i \neq \mathbf{0}$?

Because of some beautiful coincidences: 1) it turns out that $\mathbb{P}[\mathbf{z}^i | \mathbf{x}^i, \sigma, W]$ is a Gaussian, 2) for Gaussians, mean is mode which means that 3) $\mu_{\mathbf{z}}^i$ is the MLE solution to the problem $\mathbf{x}^i = W\mathbf{z}^i + \epsilon^i$ which is indeed a vector least squares regression problem

Why does the expression $\mu_{\mathbf{z}}^i = (W^T W + \sigma^2 \cdot I_k)^{-1} W^T \mathbf{x}^i$ look like least squares?

This gives us one of the alternating updates, let's derive the other



PPCA – Alternating Optimization

24

$$\sum_{i=1}^n \ln \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] = \frac{1}{2\sigma^2} \sum_{i=1}^n \|\mathbf{x}^i - W\mathbf{z}^i\|_2^2 + dn \ln \sigma$$

Thus, if \mathbf{z}^i are fixed, we can obtain W, σ using first order optimality

$$\arg \min_W \sum_{i=1}^n \|\mathbf{x}^i - W\mathbf{z}^i\|_2^2 = \arg \min_W \text{tr}(W^\top W A) - 2 \text{tr}(W^\top B)$$

where $A = \sum_{i=1}^n \mathbf{z}^i (\mathbf{z}^i)^\top$ and $B = \sum_{i=1}^n \mathbf{x}^i (\mathbf{z}^i)^\top$

See [BIS] Chap 12 for detailed derivations - check that dimensionalities match

Apply first order optimality to get $\hat{W}_{\text{MLE}} = BA^{-1}$

Once W is known, σ can also be found using first order optimality

$$\hat{\sigma}_{\text{MLE}} = \sqrt{\frac{1}{dn} \sum_{i=1}^n \|\mathbf{x}^i - \hat{W}_{\text{MLE}} \mathbf{z}^i\|_2^2}$$




ALTERNATING OPTIMIZATION

1. Initialize W
2. For $t = 0, 1, 2, \dots$
 1. Update \mathbf{z}^i fixing W, σ
 1. Let $\mathbf{z}^i = (W^\top W + \sigma^2 \cdot I_k)^{-1} W^\top \mathbf{x}^i$, for $i \in [n]$
 2. Update W, σ fixing \mathbf{z}^i
 1. Calculate $A = \sum_{i=1}^n \mathbf{z}^i (\mathbf{z}^i)^\top$ and $B = \sum_{i=1}^n \mathbf{x}^i (\mathbf{z}^i)^\top$
 2. Update $W = BA^{-1}$
 3. Calculate $V = \sum_{i=1}^n \|\mathbf{x}^i - W\mathbf{z}^i\|_2^2$
 4. Update $\sigma = \sqrt{V/dn}$



PPCA


Total of $\mathcal{O}(k^2(d + n) + kdn)$ time taken per iteration. In contrast, PCA using power + peeling method takes only $\mathcal{O}(kdn)$ time if we do power updates as $A\mathbf{v} = X^\top(X\mathbf{v})$ which takes $\mathcal{O}(dn)$ time per power step



1. Initialize W
2. For $t = 0, 1, 2, \dots$
 1. Update \mathbf{z}^i fixing W, σ
 1. Let $\mathbf{z}^i = (W^\top W + \sigma^2)^{-1} X^\top \mathbf{x}^i$
 2. Update W, σ fixing \mathbf{z}^i
 1. Calculate $A = \sum_{i=1}^n \mathbf{z}^i (\mathbf{z}^i)^\top$
 2. Update $W = BA^{-1}$
 3. Calculate $V = \sum_{i=1}^n \|\mathbf{x}^i - W\mathbf{z}^i\|^2$
 4. Update $\sigma = \sqrt{V/dn}$

$\mathcal{O}(k^3)$ time to calculate the inverse term and $\mathcal{O}(kdn + k^2d)$ time to calculate all \mathbf{z}^i terms afterward

In practice, PCA is usually faster than AltOpt PPCA – no inverses required to solve PCA, just simple mat-vec multiplication steps ☺



$\mathcal{O}(k^2n + k^3)$ time for A^{-1} , $\mathcal{O}(dkn)$ time for B and finally $\mathcal{O}(k^2d)$ time for W



PPCA – Expectation Maximization

27

We wish to get $\arg \max_{W, \sigma} \sum_{i=1}^n \ln \left(\int_{\mathbb{R}^k} \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] d\mathbf{z}^i \right)$

As before, given a model $\hat{W}, \hat{\sigma}$, let $q(\mathbf{z}^i) = \mathbb{P}[\mathbf{z}^i \mid \mathbf{x}^i, \hat{\sigma}, \hat{W}]$ and lower bound the integral using Jensen's inequality (the term C does not depend on W, σ)

$$\ln \left(\int_{\mathbb{R}^k} \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] d\mathbf{z}^i \right) \geq \int_{\mathbb{R}^k} q(\mathbf{z}^i) \cdot \ln \mathbb{P}[\mathbf{x}^i, \mathbf{z}^i \mid \sigma, W] d\mathbf{z}^i + C$$

Some simple (but non-trivial) calculations show that the resulting EM algorithm looks very similar to the AltOpt with a few simple changes

Replace \mathbf{z}^i with $\boldsymbol{\zeta}^i \triangleq \mathbb{E}[\mathbf{z}^i] = \int_{\mathbb{R}^k} q(\mathbf{z}^i) \cdot \mathbf{z}^i d\mathbf{z}^i = (W^\top W + \sigma^2 \cdot I_k)^{-1} W^\top \mathbf{x}^i$

This is the same as we used in AltOpt since for Gaussian, mode is the same as mean

$$\begin{aligned} \text{Replace } \mathbf{z}^i (\mathbf{z}^i)^\top \text{ with } Z^i &\triangleq \mathbb{E} \left[\mathbf{z}^i (\mathbf{z}^i)^\top \right] = \int_{\mathbb{R}^k} q(\mathbf{z}^i) \cdot \mathbf{z}^i (\mathbf{z}^i)^\top d\mathbf{z}^i \\ &= \boldsymbol{\zeta}^i (\boldsymbol{\zeta}^i)^\top + \sigma^2 \cdot (W^\top W + \sigma^2 \cdot I_k)^{-1} \end{aligned}$$

Rest of the algorithm (can be shown to) remain the same (see [BIS] Chap 12)




EXPECTATION MAXIMIZATION

1. Initialize W
2. For $t = 0, 1, 2, \dots$
 1. Update $\boldsymbol{\zeta}^i, Z^i$ fixing W, σ
 1. Let $\boldsymbol{\zeta}^i = (W^\top W + \sigma^2 \cdot I_k)^{-1} W^\top \mathbf{x}^i$, for $i \in [n]$
 2. Let $Z^i = \boldsymbol{\zeta}^i (\boldsymbol{\zeta}^i)^\top + \sigma^2 \cdot (W^\top W + \sigma^2 \cdot I_k)^{-1}$
 2. Update W, σ fixing $\boldsymbol{\zeta}^i, Z^i$
 1. Update $W = \left(\sum_{i=1}^n \mathbf{x}^i (\boldsymbol{\zeta}^i)^\top \right) \left(\sum_{i=1}^n Z^i \right)^{-1}$
 2. Calculate $V = \sum_{i=1}^n \|\mathbf{x}^i - W \boldsymbol{\zeta}^i\|_2^2$
 3. Update $\sigma = \sqrt{V/dn}$



PPCA – Expectation Maximization

Time complexity of PPCA using EM roughly same as that of PPCA using AltOpt.



EXPECTATION MAXIMIZATION

1. Initialize W
2. For $t = 0, 1, 2, \dots$
 1. Update $\boldsymbol{\zeta}^i, Z^i$ fixing W, σ
 1. Let $\boldsymbol{\zeta}^i = (W^\top W + \sigma^2 \cdot I_k)^{-1} W^\top \mathbf{x}^i$, for $i \in [n]$
 2. Let $Z^i = \boldsymbol{\zeta}^i (\boldsymbol{\zeta}^i)^\top + \sigma^2 \cdot (W^\top W + \sigma^2 \cdot I_k)^{-1}$
 2. Update W, σ fixing $\boldsymbol{\zeta}^i, Z^i$
 1. Update $W = \left(\sum_{i=1}^n \mathbf{x}^i (\boldsymbol{\zeta}^i)^\top \right) \left(\sum_{i=1}^n Z^i \right)^{-1}$
 2. Calculate $V = \sum_{i=1}^n \|\mathbf{x}^i - W \boldsymbol{\zeta}^i\|_2^2$
 3. Update $\sigma = \sqrt{V/dn}$



Wrapping Up Linear Models

30

Have studied how linear models can be used to do classification (binary/multi), regression, clustering, dimensionality reduction

We looked at several techniques to do so

Function Approximation Method: *define a loss function over the output (and optionally a regularizer over model parameters) and minimize (regularized) loss to learn a good model*

Probabilistic Method: *define a likelihood over the output (and optionally a prior over model parameters) and use MLE/MAP to learn a good model*

Bayesian Method: *use likelihood and prior to learn a posterior distribution over all (possibly infinite) models. At test time, aggregate responses from each one of these models, weighted by their posterior prob. to make final prediction*

Probabilistic and Bayesian methods usually come with a “generative story” specified by likelihood+prior dists. of how we assume data was generated



Non Linear Models

31

Agenda for the next few weeks (after mid-sem recess)

- Kernel Learning

- Deep Learning

Let us take half-a-step towards learning non-linear models today



Generalized Linear Models (GLM/GLiM) 32

Linear regression approximates label y^i as a linear func of features \mathbf{x}^i
... and we try to learn a model \mathbf{w} such that $y^i \approx \mathbf{w}^\top \mathbf{x}^i$

Thus, output is modelled as a linear function of features – not very flexible or powerful. Will learn techniques that model output as non-linear func of input

Generalized linear models are “somewhat” linear models

In fact, we have already seen GLMs in action – logistic reg, linear reg

In logistic regression, the label $y^i \in \{0,1\}$ (or $y^i \in \{-1,1\}$). However, $\mathbf{w}^\top \mathbf{x}^i$ doesn't give us y^i directly, rather it gives us $\mathbb{P}[y^i = 1]$. We pass $\mathbf{w}^\top \mathbf{x}^i$ through a (non-linear) thresholding function to obtain the final (binary) label

GLMs generalize this behaviour very nicely to allow us to make predictions on discrete/continuous labels even when $\mathbf{w}^\top \mathbf{x}^i$ does not directly give us y^i

Thus, GLMs are only “just” non-linear 😊 – they learn linear scores but then threshold (or apply other non-linear wrappers) to predict interesting labels

Generalized Linear Models

33

GLMs assume likelihood models of a special kind (usually $\eta = \mathbf{w}^\top \mathbf{x}$)

$$\mathbb{P}[y \mid \eta] \propto \exp(y \cdot \eta - A(\eta) - h(y))$$

*Called **canonical exponential family distributions** (details in CS772 etc)*

The normalization “constant” is sometimes absorbed into $A(\cdot)$

Sometimes, we include a “dispersion” parameter to control variance i.e.

$$\mathbb{P}[y \mid \eta, \phi] \propto \exp((y \cdot \eta - A(\eta) - h(y))/\phi)$$

Example: Gaussian ($A(\eta) = \eta^2/2, h(y) = y^2/2, \eta = \mathbf{w}^\top \mathbf{x}$)

$$\mathbb{P}[y \mid \mathbf{x}, \mathbf{w}, \sigma^2] \propto \exp\left(\left(y \cdot \mathbf{w}^\top \mathbf{x} - \frac{(\mathbf{w}^\top \mathbf{x})^2}{2} - \frac{y^2}{2}\right)/\sigma^2\right) \propto \exp\left(-\frac{(y - \mathbf{w}^\top \mathbf{x})^2}{2\sigma^2}\right)$$

Example: Bernoulli ($A(\eta) = \ln(1 + \exp(\eta)), h(y) = 0, \eta = \mathbf{w}^\top \mathbf{x}$)

$$\mathbb{P}[y = 1 \mid \mathbf{x}, \mathbf{w}] = \exp(\mathbf{w}^\top \mathbf{x} - \ln(1 + \exp(\mathbf{w}^\top \mathbf{x}))) = 1/(1 + \exp(-\mathbf{w}^\top \mathbf{x}))$$

$$\mathbb{P}[y = 0 \mid \mathbf{x}, \mathbf{w}] = \exp(-\ln(1 + \exp(\mathbf{w}^\top \mathbf{x}))) = 1/(1 + \exp(\mathbf{w}^\top \mathbf{x}))$$

Generalized Linear Models

34

Example: Poisson ($A(\eta) = \exp(\eta)$, $h(y) = \ln(y!)$, $\eta = \mathbf{w}^\top \mathbf{x}$)

The Poisson distribution can model labels that take values in \mathbb{N} . Typically used if we want to predict counts e.g. number of students graduating etc

$$\mathbb{P}[y = k \mid \mathbf{x}, \mathbf{w}] = \frac{1}{k!} \exp(k \cdot \mathbf{w}^\top \mathbf{x} - \exp(\mathbf{w}^\top \mathbf{x})) = \frac{\lambda^k e^{-\lambda}}{k!} \text{ with } \lambda = \exp(\mathbf{w}^\top \mathbf{x})$$

Example: Gamma ($A(\eta) = \ln(-\eta)$, $h(y) = (1 - \phi^{-1}) \ln(y)$, $\eta = -\exp(\mathbf{w}^\top \mathbf{x})$)

Almost like a continuous Poisson – can model labels that take only non-negative values e.g. time before patient visits again


Note: $\eta = -\exp(\mathbf{w}^\top \mathbf{x})$ ensures $\eta < 0$ o/w gamma dist doesn't make sense

$$\mathbb{P}[y \mid \mathbf{x}, \mathbf{w}, \phi] \propto \exp((-y \cdot \exp(\mathbf{w}^\top \mathbf{x}) + \mathbf{w}^\top \mathbf{x} + (1 - \phi) \ln y) / \phi)$$

Example: von-Mises likelihood $\mathbb{P}[y \mid \mathbf{x}, \mathbf{w}, \kappa] \propto \exp(\kappa \cos(y - \mathbf{w}^\top \mathbf{x}))$

A Gaussian on a circle/finite interval – usually used to predict angles $\in [0, 2\pi]$

A “non-canonical” exponential family member – form of expression different



Mostly, we have $\mathbb{P}[y \mid \mathbf{x}, \mathbf{w}] \propto \exp(y \cdot \mathbf{w}^\top \mathbf{x} - A(\mathbf{w}^\top \mathbf{x}) - h(y))$

Interesting fact: *always have the mean predicted label $\mathbb{E}[y \mid \mathbf{x}, \mathbf{w}] = A'(\eta)$*

Gaussian: $A'(\eta) = \eta = \mathbf{w}^\top \mathbf{x}$

Bernoulli: $A'(\eta) = 1/(1 + \exp(-\eta)) = \sigma(-\mathbf{w}^\top \mathbf{x})$ ($\sigma \equiv$ sigmoid function)

Poisson: $A'(\eta) = \exp(\eta) = \exp(\mathbf{w}^\top \mathbf{x})$

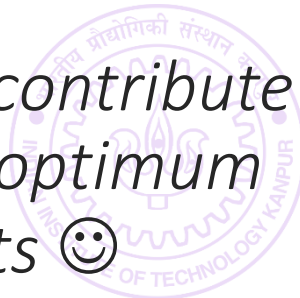
Gamma: $A'(\eta) = -1/\eta = \exp(-\mathbf{w}^\top \mathbf{x})$ (recall in this case $\eta = -\exp(\mathbf{w}^\top \mathbf{x})$)

Solving for MLE becomes simple (\mathcal{C} does not depend on \mathbf{w}) via (S)GD

$$\mathcal{L}(\mathbf{w}) = -\ln \mathbb{P}[\mathbf{y} \mid X, \mathbf{w}] = \sum_{i=1}^n A(\mathbf{w}^\top \mathbf{x}^i) - y^i \cdot \mathbf{w}^\top \mathbf{x}^i + \mathcal{C}$$

Thus, we have $\nabla \mathcal{L}(\mathbf{w}) = \sum_{i=1}^n (A'(\mathbf{w}^\top \mathbf{x}^i) - y^i) \cdot \mathbf{x}^i$

Note that if for some point $A'(\mathbf{w}^\top \mathbf{x}^i) = y^i$ then that point does not contribute to the gradient. One way to get zero gradient (which will be a global optimum since $\mathcal{L}(\mathbf{w})$ is mostly convex) is to ensure $A'(\mathbf{w}^\top \mathbf{x}^i) = y^i$ for all points 😊



MLE with GLMs

36

Mostly, we have $\mathbb{P}[y \mid \mathbf{x}, \mathbf{w}] \propto \exp(y \cdot \mathbf{w}^\top \mathbf{x} - A(\mathbf{w}^\top \mathbf{x}) - h(y))$

Interesting fact: *always have the mean predicted label $\mathbb{E}[y \mid \mathbf{x}, \mathbf{w}] = A'(\eta)$*

Gaussian: $A'(\eta) = \eta = \mathbf{w}^\top \mathbf{x}$

Bernoulli: $A'(\eta) = 1/(1 + \exp(-\eta))$

Poisson: $A'(\eta) = \exp(\eta) = \exp(\mathbf{w}^\top \mathbf{x})$

Gamma: $A'(\eta) = -1/\eta = \exp(-\eta)$

Slightly different (but similar) updates required for gamma since we do not have $\eta = \mathbf{w}^\top \mathbf{x}$ but rather $\eta = -\exp(\mathbf{w}^\top \mathbf{x})$ in that case

Solving for MLE becomes simple (\mathcal{C} does not depend on \mathbf{w}) via (S)GD

$$\mathcal{L}(\mathbf{w}) = -\ln \mathbb{P}[\mathbf{y} \mid X, \mathbf{w}] = \sum_{i=1}^n A(\mathbf{w}^\top \mathbf{x}^i) - y^i \cdot \mathbf{w}^\top \mathbf{x}^i + \mathcal{C}$$

$$\text{Thus, we have } \nabla \mathcal{L}(\mathbf{w}) = \sum_{i=1}^n (A'(\mathbf{w}^\top \mathbf{x}^i) - y^i) \cdot \mathbf{x}^i$$

Note that if for some point $A'(\mathbf{w}^\top \mathbf{x}^i) = y^i$ then that point does not contribute to the gradient. One way to get zero gradient (which will be a global optimum since $\mathcal{L}(\mathbf{w})$ is mostly convex) is to ensure $A'(\mathbf{w}^\top \mathbf{x}^i) = y^i$ for all points 😊

